

# Linked lists in C

Mo Wang

Spring 2026

## Introduction

This report investigates and delves into performance and behavior of a linked list in C. While linked list offers flexibility by carrying an address of next element of a list, it has some performance bottlenecks in terms of random access and cache locality. The primary goal of this report is to analyze the behavior and key attribute of a linked list and compare these features with arrays.

## Linked list

A linked list stores a sequence of elements by holding its first element. Each element in a linked list is stored in a cell with an element value and a pointer with the address of the next cell. A null pointer in a linked list cell denotes the last cell in the list, while a null pointer as the first element denotes the absence of elements in the list.

```
typedef struct cell {
    int value;
    struct cell *tail;
} cell;

typedef struct linked {
    cell *first;
} linked;
```

The linked list can be created by allocating a memory for the linked list and successively allocating other cells of the list, where one cell address is stored in the previous one's tail.

A linked list is created by first allocating memory for the list structure itself. Each cell is then successively allocated with each cell's tail pointer storing the address of the next cell in the sequence, hence linking all the cells together. The linked list structure stores a pointer to the first cell as

an entry point to the data structure- From a memory perspective, these cells are scattered in dynamic memory, which provides flexibility at the expense of high performance overhead of accessing and searching operations.

```
linked *linked_create() {
    linked *new = (linked*)malloc(sizeof(linked));
    new->first = NULL;
    return new;
}
void linked_free(linked *lnk) {
    cell *nxt = lnk->first;
    while (nxt != NULL) {
        cell *tmp = nxt->tail;
        free(nxt);
        nxt = tmp;
    }
    free(lnk);
}
```

### Adding an element

To insert an element at the beginning of a linked list, a new cell is first allocated. The value of the new cell is assigned to the desired element. Its tail pointer is then assigned to point to the current first cell of the list. Finally, the first pointer of the list is updated to point to this new cell, effectively making it the first element of the list. Since the operation requires only one memory allocation and three pointer updates, the time complexity is constant  $O(1)$ .

```
void linked_add(linked *lnk, int item) {
    cell *new = (cell*) malloc(sizeof(cell));
    new->value = item;
    new->tail = lnk->first;
    lnk->first = new;
}
```

### Length of linked list

Because this implemented linked list structure only stores a pointer to first element cell without other necessary metadata, such as the size of the list, calculating the length of the linked list could only be accomplished by traversing the linked list through all the cells. A length counter is held and increments when moving to another cell through the previous cell's tail pointer during traversal. As all cells must be traversed, the complexity of the operation in time is linear to the size of the list  $O(n)$ .

```

unsigned int linked_length(linked *lnk){
    unsigned int length = 0;
    cell* nxt = lnk->first;
    while (nxt != NULL){
        nxt = nxt->tail;
        length++;
    }
    return length;
}

```

## Finding an element

Finding an element requires traversing the linked list from the first element cell to the last one, since the list can only be efficiently accessed sequentially. In this case, the value of each element cell is compared with the search value. The traversal search stops as the element is found. This results in linear time complexity  $O(n)$ , since each cell has to be traversed in the worst case, while only one cell is traversed in the best case. The operation performs on average  $n/2$  traversal in terms of possibility expectation.

```

bool linked_find(linked *lnk, int item){
    cell *nxt = lnk->first;
    while (nxt != NULL && nxt->value != item){
        nxt = nxt->tail;
    }
    return nxt != NULL;
}

```

## Removing an element

An element can be removed by finding the element cell by value and therefore unlink it from adjacent cells. The operation is split into two cases: either the first element is targeted for deletion, or the deleted element is in the middle of the list. Removing the first element can be accomplished by updating the first pointer of linked list to the second cell while freeing the first cell. Meanwhile, removing an element in the middle of the array requires finding the element by tracking its previous cell. As the target element cell is found, the previous cell re-links its tail to the element after the target one, while freeing the target afterwards. Although pointer updates and free operations are done in constant time, the target element cell must be located by finding, resulting in linear time complexity  $O(n)$ , as described earlier.

```

void linked_remove(linked *lnk, int item){
    cell *nxt = lnk->first;
    if (nxt == NULL) return;

```

```

if (nxt->value == item){
    cell *after = nxt->tail;
    free(nxt);
    lnk->first = after;
}
else{
    while (nxt->tail != NULL && nxt->tail->value != item){
        nxt = nxt->tail;
    }
    if (nxt->tail != NULL){
        cell *after = nxt->tail->tail;
        free(nxt->tail);
        nxt->tail = after;
    }
}
}

```

## Appending a linked list

Two linked lists can be appended by updating the tail pointer of the last cell of the first linked list `a` to the first cell of the other one. Afterwards, the other linked list clears its cells by pointing the first element cell to null. Due to the absence of a last element cell pointer in the linked list structure, accessing the last cell of the first linked list can only be performed by traversing the first array to the last element cell, which has a time complexity linear to the first array size.

As a result, the overall time complexity of appending one linked list to another in this scenario is  $O(n)$ , where  $n$  is the number of elements in the first linked list. This contrasts with an implementation that maintains a direct tail pointer for each linked list, which would allow appending in  $O(1)$  constant time, since the last element could be accessed immediately without traversal.

```

void linked_append(linked *a, linked *b) {
    cell *nxt = a->first;
    cell *prv = NULL;
    while(nxt != NULL) {
        prv = nxt;
        nxt = nxt->tail;
    }
    if (prv == NULL){
        a->first = b->first;
    }
    else{
        prv->tail = b->first;
    }
}

```

```

    }
    b->first = NULL;
}

```

### Benchmark of append operation

The benchmark results also conforms the linear time complexity growth of linked list append with a growing left linked list during append operation, alongside with theoretical performance analysis. The linear growth behavior  $O(n)$  is shown in Figure 1, as the left linked list varies in size, while the right linked list is kept in constant size of 8192 elements, with execution time data shown in Table 1.

Size	1024	2048	4196	8192	16384	32768	65535
Time (ns)	$3.7 \times 10^3$	$6.6 \times 10^3$	$1.2 \times 10^4$	$2.4 \times 10^4$	$4.7 \times 10^4$	$9.4 \times 10^4$	$1.9 \times 10^5$

Table 1: Minimum time per loop for append benchmark (transposed)

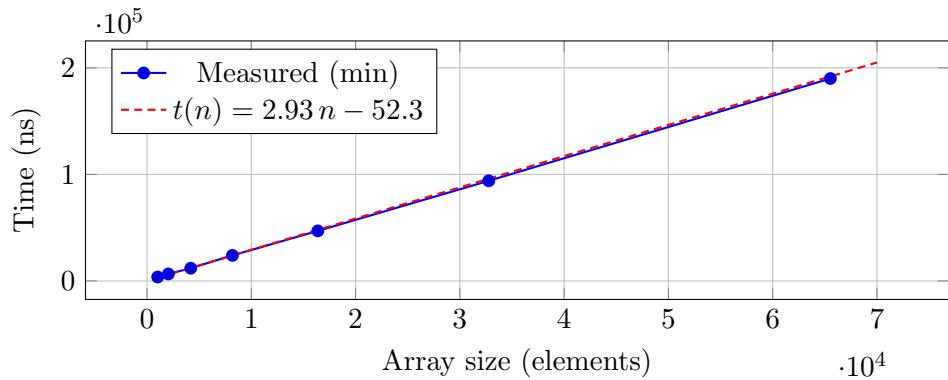


Figure 1: Minimum measured time per loop with fitted linear regression

Switching the roles between the two linked list, where the left one is kept fixed while the other one growing in size results into a constant time growth behavior in Figure 2, with execution time data in Table 2. This also conforms with the theoretical performance analysis, since the last element cell of the left linked list must be traversed before constant time operation of pointer updates can be performed.

Size	1024	2048	4196	8192	16384	32768	65535
Time (ns)	$2.4 \times 10^4$						

Table 2: Minimum time per loop (transposed)

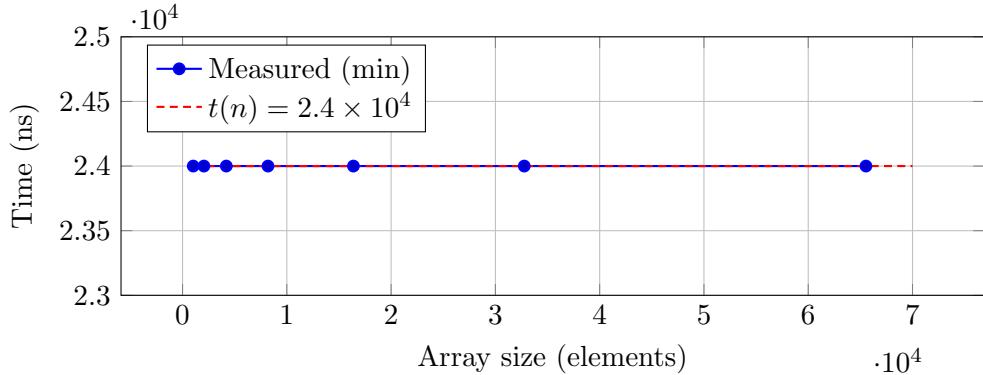


Figure 2: Minimum measured time per loop with fitted constant regression

Assuming the left linked list contains  $n$  elements, while the right one  $m$  elements, the time complexity grows linearly with the left one. Thus, the time complexity can be concluded as  $O(n)$  and does not depend on the right one  $m$ .

## Comparison with an array

By comparing performance of a linked list to an array in terms of appending two lists, linked list slightly outperforms arrays. The main reason is that arrays are fixed in size and must be copied into another array when resizing. Appending two arrays requires resizing the left array and copy the elements of the right array into the left one with sufficient space. In this case, the cost of appending two arrays is copying all the elements in both arrays. Assuming the left one has  $n$  elements and the right one  $m$  elements, the total time complexity is  $O(n + m)$ . Compared to linked list, the time complexity is only  $O(n)$ . In Figure 3, linear growth can be visualized with execution time data from Table 3.

Size	1024	2048	4196	8192	16384	32768	65535
Time (ns)	$2.0 \times 10^4$	$2.2 \times 10^4$	$2.7 \times 10^4$	$4.5 \times 10^4$	$6.2 \times 10^4$	$8.3 \times 10^4$	$1.5 \times 10^5$

Table 3: Minimum time per loop vs. input size (transposed)

## Stack implementation using linked list

A stack can also be implemented using a linked list. The elements are pushed to the first position in the linked list, which can thereafter be popped by removing the first element of the list and returning its value. The push

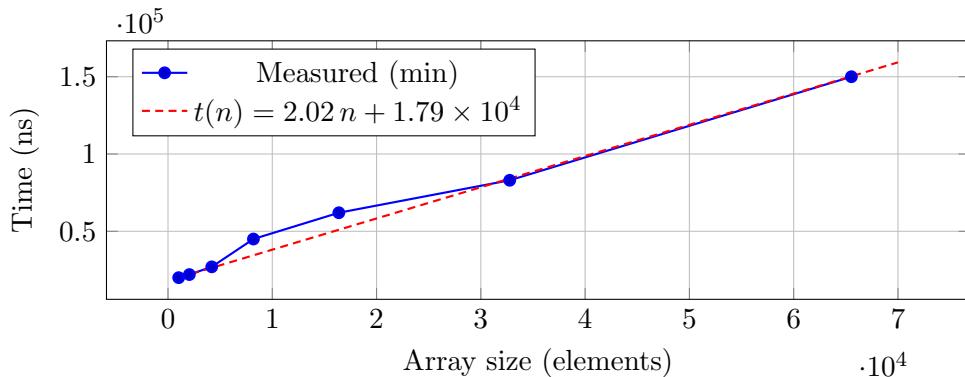


Figure 3: Minimum time per loop with fitted linear trend

operation is implemented exactly the same as adding an element into the first position of the linked list. However, the pop operation requires first removing the first element from the linked list stack, by updating the pointer to first element cell to point to the second element, while de-allocating the first element.

```

void linked_stack_push(linked_stack *lnk, int item) {
    // same as inserting element into first position in linked list
}

int linked_stack_pop(linked_stack *lnk) {
    if (lnk->first == NULL) {
        return 0;
    }
    // Remove first element
    linked_stack_cell *first_cell = lnk->first;
    int value = first_cell->value;
    lnk->first = lnk->first->tail;
    free(first_cell);
    return value;
}

```

Although a linked-list stack grows flexibly without resizing, it is slower in practice. Both array-based and linked-list stacks have amortized  $O(1)$  push/pop time, but linked lists suffer from pointer chasing, cache misses, and allocation overhead. Arrays use contiguous memory with cheap reads and writes, so they generally outperform linked lists despite identical asymptotic complexity.

## Conclusion

A linked list in C offers structural flexibility with inexpensive insertions and deletions, while arrays benefit from contiguous memory and superior cache behavior. Searching in a linked list remains linear time  $O(n)$ , whereas insertion or deletion is  $O(1)$  when the target cell is known.

The append benchmark shows that, without a tail pointer, a linked list grows with linear time complexity  $O(n)$  because the last element must be located before linking. In contrast, appending arrays incurs a dominant copying cost of  $O(n+m)$  when resizing is required, while the pointer update itself is constant time  $O(1)$ . A tail pointer would reduce linked-list append to constant time as well.

Aspect	Linked list	Array
Insertion / deletion	$O(1)$ if the cell pointer is known; supports flexible structural modifications	Expensive shifts for middle insert/delete; fixed-size layout (unless dynamically resized)
Search	Requires full traversal; $O(n)$	Fast random access; $O(1)$ per access, but search is $O(n)$ unless sorted or indexed
Append operation	$O(n)$ without a tail pointer (must traverse to last element); becomes $O(1)$ if a tail pointer is maintained	Requires copying elements into a resized array; $O(n+m)$ when appending arrays of size $n$ and $m$
Memory layout	Elements scattered across the heap; poor cache locality	Contiguous memory; excellent cache performance and predictable access patterns
Practical performance	Flexible but often slower due to pointer chasing and heap allocation	Typically faster in real workloads due to locality and cheap read/write operations

Table 4: Comparison of linked list and array operations

Overall, linked lists trade memory locality for flexibility, whereas arrays achieve faster practical performance thanks to contiguous storage and predictable access patterns.